

# eBPF for Python Troubleshooting


2024-11-23




**PYCHINA**

Python 中国社区



- Cilium datapath tech lead 
- Started my career from Python
- Debug for fun



## Let's debug, in incident mode

1. CVE-2024-27351
2. Django ticket #27690
3. BentoML issue #4760

### Don'ts

1. No time to read the whole code
2. No idea about impl details
3. Non-disruptive required

### Dos

1. Observe symptoms
2. Collect metrics
3. Non-disruptive debuggers

### Notes:

- Containerized CPython 3.10 (docker run python:3.10)
- GNU/Linux 6.8.0 (Ubuntu 22.04.1)
- x86\_64, Little Endian
- CONFIG\_BPF\*=y, CONFIG\_KPROBE\*=y (cat /boot/config-\$(uname -r))
- bpftrace v0.20.3

CVE-2024-27351

# Symptoms

1. Django HTTP server suddenly doesn't serve (`curl localhost:8000` hangs/timeout/slow)
2. Django Python process consumes CPU (`top`)

Random guesses:

- App bugs: infinite loop?
- Bad optimization: deep recursion?
- No bug: too many requests?

Idea: Find the Django code responsible for CPU spike, likely that's the cause of slow response.

# perf

```
$ perf record -F99 -g -p $PID --call-graph dwarf,16384
$ perf report -g
```

```
# Children  Self  Command  Shared Object  Symbol
# .....  .....  .....  .....  .....
#
 96.43%    0.00%  python  [unknown]      [k] 0xffffffffffffffff
      |
      |--0xffffffffffffffff
      |
      |--74.02%--_PyObject_FastCallDictTstate
      |         _PyEval_Vector
      |         _PyEval_EvalFrameDefault
      |         _PyEval_Vector
      |         _PyEval_EvalFrameDefault
      |         _PyEval_Vector
      |         _PyEval_EvalFrameDefault
      |         _PyObject_MakeTpCall
      |         slot_tp_call
      |         _PyObject_Call_Prepend
      |         _PyEval_EvalFrameDefault
      |         cfunction_vectorcall_FASTCALL_KEYWORDS_METHOD
      |         _sre_SRE_Pattern_search
      |         sre_search
      |
      |--71.02%--sre_ucs1_match
      |
      |--3.00%--sre_ucs1_match.cold
```

ChatGPT, what's sre\_ucs1\_match in Python?



: “an internal function in re module”

Looks like **ReDoS** (Regular expression Denial of Service)! But which Python code called what regex?

# Dark side of perf (In the context of CPython)

1. Output C functions rather than Python functions
2. Requirements in how CPython is built
  - a. `-fno-omit-frame-pointer`: with or without frame pointer (affecting algorithm and overhead)
  - b. `-s`: with or without symbols (affecting address-symbol mapping)

# python3.12+ -X perf

```
$ python3.12 -X perf manage.py runserver  
$ perf record -F99 -g -p $PID --call-graph dwarf,16384  
$ perf report -g
```

```
96.90%   82.39%  python  libpython3.12.so.1.0      [...] sre_ucs1_match  
|  
|--78.24%--py::Truncator._truncate_html:/usr/local/lib/python3.12/site-packages/django/utils/text.py  
|   _PyEval_EvalFrameDefault  
|   PyObject_Vectorcall  
|   cfunction_vectorcall_FASTCALL_KEYWORDS_METHOD  
|   _sre_SRE_Pattern_search  
|   sre_ucs1_search  
|   sre_ucs1_match
```

It's `Truncator()._truncate_html()` from `/.../django/utils/text.py` that finally calls `sre_ucs1_match`.

But who called the `_truncate_html()`? Can we see the whole Python call chain?



# Dark side of -X perf

1. Official images not built with `-fno-omit-frame-pointer`
  - a. Can't get full Python function call chain
  - b. `python3.13+ -X perf_jit` (kernel  $\geq 6.7.3$ , higher overhead)
2. Performance overhead
  - a. Requests/second: 142.63  $\rightarrow$  123.90, -13% (`ab -n10000 localhost:8000/`)
  - b. Extreme case: -56% (recursive fib)
3. Require process restart to specify `-X perf`
  - a. Prod env doesn't allow
  - b. Issues may be gone after restart
  - c. Why don't `python -m cProfile`?

# py-spy

\$ py-spy record -p \$PID

start_django (django/utils/autoreload.py:660)	call (django/core/handlers/wsgi.py:124)	
run (django/utils/autoreload.py:344)	get_response (django/core/handlers/base.py:140)	< -
run_loop (django/utils/autoreload.py:350)	inner (django/core/handlers/exception.py:55)	
tick (django/utils/autoreload.py:406)	call (django/utils/deprecation.py:134)	
	inner (django/core/handlers/exception.py:55)	
	call (django/utils/deprecation.py:134)	
	inner (django/core/handlers/exception.py:55)	
	call (django/utils/deprecation.py:134)	
	inner (django/core/handlers/exception.py:55)	
	call (django/utils/deprecation.py:134)	
	inner (django/core/handlers/exception.py:55)	
	call (django/utils/deprecation.py:134)	
	inner (django/core/handlers/exception.py:55)	
	call (django/utils/deprecation.py:134)	
	inner (django/core/handlers/exception.py:55)	
	get_response (django/core/handlers/base.py:197)	
	view_wrapper (django/views/decorators/csrf.py:65)	
	index (polls/views.py:10)	< -
	words (django/utils/text.py:148)	< -
	truncate_html (django/utils/text.py:200)	< -

```
-> get_response[django]
-> index (polls/views.py:10)
-> words (django/utils/text.py:148)
-> _truncate_html (django/utils/text.py:200)
```

Root cause: `index()` in `polls/views.py:10` finally called `_truncate_html()` in Django, which has a regex search that triggered **ReDoS**.

# Dark side of py-spy

## 1. Can't distinguish off-CPU from on-CPU

- a. Misleading plateaus

## 2. Performance overhead

- a. Requests/second: 170.94 → 134.11, -21% (ab -n10000 localhost:8000/)
- b. py-spy **--nonblocking**: 170.94 → 164.83, -3% (but may lead to inaccurate results)

	Disruptive	Requirements in Python build flags	Overhead	Accuracy
<code>perf</code>	No	Yes	Low (sampling)	No Python stack
<code>perf + python -X perf</code>	Yes	Yes	Moderate (non-sampling)	One layer of Python stack
<code>py-spy</code>	No	No	High (sampling)	Possible misleading result

# eBPF: speedrun

Like a breakpoint, but more flexible, more efficient, more powerful

```
$ gdb -p 88991
(gdb) break sre_ucs1_match
Breakpoint 1 at 0x7ee02c2d3a70: file ./Modules/sre_lib.h, line 550.
(gdb) c
Continuing.

Thread 3 "python" hit Breakpoint 1, sre_ucs1_match (state=0x7ee0291fa2f0, pattern=0x7ee02953852c,
toplevel=0) at ./Modules/sre_lib.h:550
(gdb) backtrace

$ bpftrace -e 'uprobe:/proc/88991/root/usr/local/lib/libpython3.10.so.1.0:sre_ucs1_match
/pid == 88991/ {print(ustack)}'
```

[1] target pid (filter)

[2] breakpoint (event hook)

[3] commands to execute at breakpoint (callback)

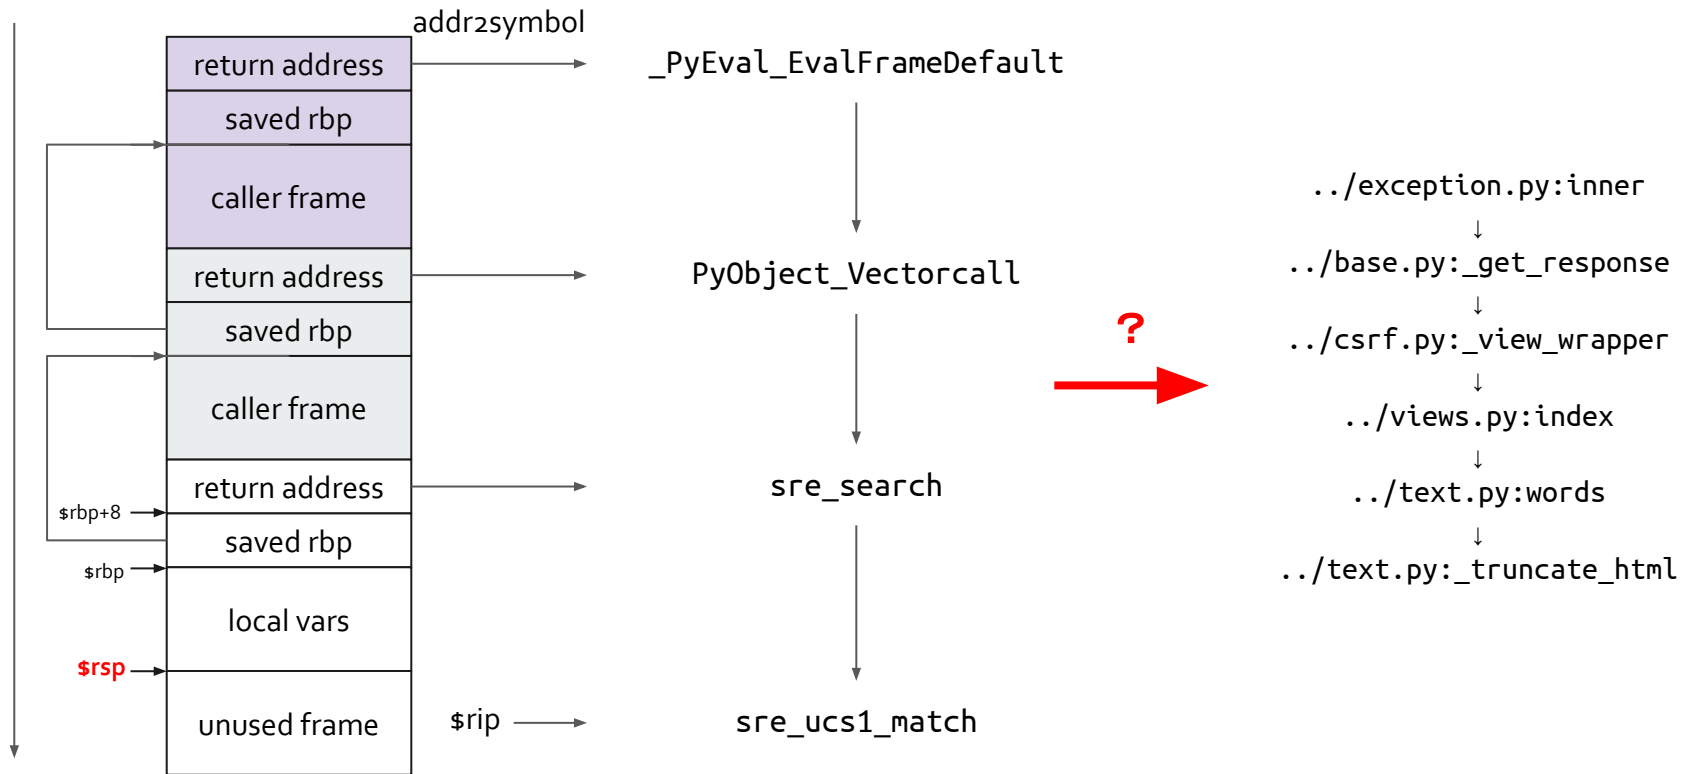
# eBPF for CPython perf: Vision

1. **Event hook**: on-CPU samples
2. **Callback**: an eBPF program to unwind stack in Python VM level (HOW?)
3. **Filter**: target Python pid

```
./.../autoreload.py:tick  
./.../autoreload.py:run_loop  
./.../autoreload.py:start_django  
./.../manage.py:main  
./.../manage.py:<module>: 1  
  
./.../selectors.py:select  
./.../socketserver.py:serve_forever  
./.../threading.py:_bootstrap_inner  
./.../threading.py:_bootstrap: 2  
  
./.../text.py:_truncate_html  
./.../text.py:words  
./.../views.py:index  
./.../deprecation.py:__call__  
./.../exception.py:inner: 260
```

This stack was seen **260** times during on-CPU sampling! Gotcha!

# Technical Hurdles



# PyFrameObject: an abstract of Python Frame

```
PyObject *_PyEval_EvalFrameDefault(PyThreadState *, PyFrameObject *, int);
```

```
(gdb) ptype/o PyFrameObject
```

```
type = struct _frame {  
    /* 0      | 24 */ PyVarObject ob_base;  
    /* 24     | 8  */ struct _frame *f_back;  
    /* 32     | 8  */ PyCodeObject *f_code;  
    /* 40     | 8  */ PyObject *f_builtins;  
    /* 48     | 8  */ PyObject *f_globals;  
    /* 56     | 8  */ PyObject *f_locals;  
    /* 64     | 8  */ PyObject **f_valuестack;  
    /* 72     | 8  */ PyObject *f_trace;  
    /* 80     | 4  */ int f_stackdepth;  
    /* 84     | 1  */ char f_trace_lines;  
    /* 85     | 1  */ char f_trace_opcodes;  
    /* XXX 2-byte hole */  
    /* 88     | 8  */ PyObject *f_gen;  
    /* 96     | 4  */ int f_lasti;  
    /* 100    | 4  */ int f_lineno;  
    /* 104    | 4  */ int f_iblock;  
    /* 108    | 1  */ PyFrameState f_state;  
    /* XXX 3-byte hole */  
    /* 112   | 240 */ PyTryBlock f_blockstack[20];  
    /* 352   | 8  */ PyObject *f_localsplus[1];  
  
    /* total size (bytes): 360 */  
};
```

Docs: frame.ob\_base has a reference to object's type (ob\_type).

Can we rely on certain characteristic to identify a PyFrameObject pointer?

Docs: frame.f\_back points to the previous stack frame (towards the caller).

Can we unwind python stack using it?

Docs: frame.f\_code holds code object being executed in this frame.

Can we read python func name and filename from it?

Can we unwind Python stack based on PyFrameObject?



# PyFrameObject->ob\_base: to find a PyFrameObject

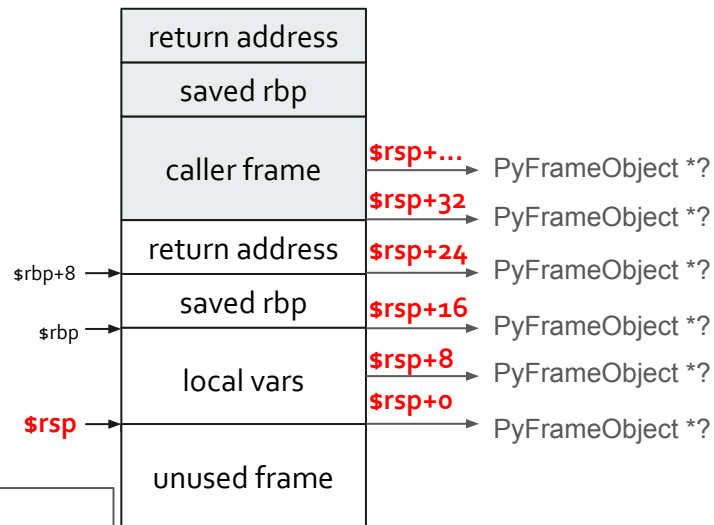
```
(gdb) bt
...
#7  _PyEval_EvalFrameDefault (tstate=<optimized out>, f=<optimized out>, throwflag=<optimized out>)
    at Python/ceval.c:4181
#8  0x00007b2ee9c0992a in _PyEval_EvalFrame (throwflag=<optimized out>, f=0x7b2ee78b72e0,
    tstate=0x60cadbf29190) at ./Include/internal/pycore_ceval.h:46
```

```
(gdb) f 8
(gdb) p f
$10 = (PyFrameObject *) 0x7b2ee78b72e0
```

```
(gdb) p f->ob_base
(gdb) ptype f->ob_base
(gdb) p f->ob_base.ob_base.ob_type->tp_name
$12 = 0x7b2ee9c93ba4 "frame"
```

```
(gdb) p ((PyFrameObject*)0x7b2ee78b72e0)->ob_base.ob_base.ob_type->tp_name
$12 = 0x7b2ee9c93ba4 "frame"
```

```
(gdb) p ((PyFrameObject*)0x7b2ee78b7233)->ob_base.ob_base.ob_type->tp_name
Cannot access memory at address 0x24
```



Conclusion: We can find a PyFrameObject by finding an stack address where

```
(PyFrameObject*)$addr->ob_base.ob_base.ob_type->tp_name == "frame"
```

# PyFrameObject->f\_back: to unwind Python stack

```
(gdb) p f
$10 = (PyFrameObject *) 0x7b2ee78b72e0
```

```
(gdb) ptype f
type = struct _frame {
...
} *
```

```
(gdb) p f->f_back
$11 = (struct _frame *) 0x60cad2f9780
```

Conclusion: We can unwind stack in Python VM level via  
**PyFrameObject->f\_back!**

```
struct PyFrameObject *f;

for (int i=0; i<20; i++) {
    f = f->f_back;
    if (!f)
        break;
}
```

# PyFrameObject->f\_code: to get Python func/file name

```
(gdb) p f
$2 = (PyFrameObject *) 0x7b2ee78b72e0
```

```
(gdb) p f->f_code
$3 = (PyCodeObject *) 0x7b2ee85f8df0
```

```
(gdb) ptype f->f_code
```

```
...
(gdb) p f->f_code->co_name
$4 = (PyObject *) 0x7b2ee85f0ef0
```

```
(gdb) x/100s f->f_code->co_name
(gdb) x/s (unsigned long long)f->f_code->co_name + 48
0x7b2ee85f0f20:      "tick"
```

```
(gdb) x/s (unsigned long long)f->f_code->co_filename + 48
0x7b2ee85dcae0:      "/usr/local/lib/python3.10/site-packages/django/utils/autoreload.py"
```

```
#define STRUCT_FOR_ASCII_STR(LITERAL) \
    struct { \
        PyASCIIObject _ascii; \
        uint8_t _data[sizeof(LITERAL)]; \
    }
```

`sizeof(PyASCIIObject) == 48`

But to be honest, you don't need to know this CPython implementation detail.

Conclusion: We can get python function name and filename via  
`PyFrameObject->f_code->{co_name,co_filename} + 48!`

# bpfftrace script

```
profile:hz:99 /pid == $1/  
{
```

0. eBPF hook: perf -F99 -p \$PID

1. Get user space register

```
$rsp = reg("sp");
```

```
$frame = (struct PyFrameObject *)0;  
$i = (uint64)0;  
$sp = (uint64)$rsp;  
while ($i <= 200) {  
    $frame = *(struct PyFrameObject**)( $sp + 8*$i);  
    if (str($frame->ob_base.ob_base.ob_type->tp_name, 5) == "frame") {  
        break;  
    }  
    $frame = (struct PyFrameObject *)0;  
    $i += 1;  
}  
  
if ($frame == 0) {  
    return;  
}
```

2. Search stack memory from `$rsp`, find the `PyFrameObject` pointer via `PyFrameObject->ob_base.ob_base.ob_type->tp_name`

3. Print py function name and filename via `PyFrameObject->f_code->{co_name,co_filename} + 48`

```
$i = 0;  
printf("\n");  
while ($i < 20) {  
    printf("%s:%s\n", $frame->f_code->co_filename->buf, $frame->f_code->co_name->buf);  
    $i += 1;  
    $frame = $frame->f_back;  
    if ($frame == 0) {  
        return;  
    }  
}
```

4. Unwind py stack via `PyFrameObject->f_back`

```
}
```

# Misc: Type Definitions (cpython310.h)

```
struct PyTypeObject {  
    char _[24];  
    char *tp_name;  
};  
struct PyObject {  
    char _[8];  
    struct PyTypeObject *ob_type;  
};  
struct PyVarObject {  
    struct PyObject ob_base;  
    char _[8];  
};
```

CPython version sensitive  
But method applies

```
struct _PyStr {  
    char _[48];  
    char buf[100];  
};  
struct PyCodeObject {  
    char _[104];  
    struct _PyStr *co_filename;  
    struct _PyStr *co_name;  
};
```

Remember `PyFrameObject->f_code->{co_name,co_filename} + 48?`  
We don't need to know why it's 48 bytes offset, just define an unused field of 48 bytes size

```
struct PyFrameObject {  
    struct PyVarObject ob_base;  
    struct PyFrameObject *f_back;  
    struct PyCodeObject *f_code;  
};
```

We don't need to define all fields

# Misc: Output (group.py)

```
.../autoreload.py:tick  
.../autoreload.py:run_loop  
.../autoreload.py:run
```

```
.../autoreload.py:tick  
.../autoreload.py:run_loop  
.../autoreload.py:run
```

```
.../text.py:_truncate_html  
.../text.py:words  
.../views.py:index
```



```
.../text.py:_truncate_html  
.../text.py:words  
.../views.py:index: 1
```

```
.../autoreload.py:tick  
.../autoreload.py:run_loop  
.../autoreload.py:run: 2
```

```
import sys  
from collections import defaultdict
```

```
stack = []  
stacks = defaultdict(int)  
try:
```

```
    for line in sys.stdin:  
        stack.append(line.strip())  
        if not line.strip():  
            stacks['\n'.join(stack)] += 1  
            stack = []
```

```
except KeyboardInterrupt:  
    pass
```

```
# output by count of stacks  
for stack, count in sorted(stacks.items(), key=lambda x: x[1]):  
    print(f'{stack.strip()}: {count}\n')
```

Read from stdin (shell pipe)

Split by paragraph

Sort by number of occurrences

# Misc: Kernel space (ubuntu2204.h)

```
#define PAGE_SIZE (1<<12)
#define KASAN_STACK_ORDER 0
#define THREAD_SIZE_ORDER (2 + KASAN_STACK_ORDER)
#define THREAD_SIZE ((uint64) (PAGE_SIZE << THREAD_SIZE_ORDER))
#define TOP_OF_KERNEL_STACK_PADDING ((uint64) 0)
```

```
profile:hz:99 /pid == $1/
{
```

```
    $rsp = reg("sp");
```

```
    // kernel mode
```

```
    $user_mode = reg("cs") & 3;
```

```
    if (!$user_mode) {
```

```
        $task = (struct task_struct *)curtask;
```

```
        $__ptr = (uint64)$task->stack;
```

```
        $__ptr += THREAD_SIZE - TOP_OF_KERNEL_STACK_PADDING;
```

```
        $pt_regs = ((struct pt_regs *)$__ptr) - 1;
```

```
        $rsp = $pt_regs->sp;
```

```
    }
```

```
    ...
}
```

This may point to kernel stack if process is trapped in kernel mode (syscall)

If in kernel mode, get user space registers from task\_struct. (Algorithm and macros may vary among different kernel versions and distros.)

Conclusion: We must get user space registers if process is in kernel space.





	Disruptive	Requirements in Python build flags	Overhead	Accuracy
perf	No	Yes	Low (sampling)	No Python stack
perf + python -X perf	Yes	Yes	Moderate (non-sampling)	One layer of Python stack
py-spy	No	No	High (sampling)	Possible misleading result
bpftrace	No	No	Super low (sampling)	100% on-CPU stacks

RPS -0% (frequency 99, python stack depth 20)

# CVE-2024-27351

## Common Vulnerabilities and Exposures

[\[Previous\]](#) [\[Index\]](#) [\[Next\]](#)

### Upstream information

[CVE-2024-27351 at MITRE](#)

#### Description

In Django 3.2 before 3.2.25, 4.2 before 4.2.11, and 5.0 before 5.0.3, the `django.utils.text.Truncator.words()` method (with `html=True`) and the `truncatewords_html` template filter are subject to a potential regular expression denial-of-service attack via a crafted string. NOTE: this issue exists because of an incomplete fix for CVE-2019-14232 and CVE-2023-43665.

Django ticket #27690

# Symptoms

1. Django DB test runs slowly (>1s) even for empty project (`time python manage.py test`)
2. Django Python process **doesn't** consume CPU (`top`)

Random guesses:

- Mutex?
- Disk IO? Network IO?
- Kernel? (scheduler? cgroups?) (unlikely in this case)

Idea: Find the **syscalls** responsible for off-CPU, likely that's the cause of slow response.

# strace

```
$ strace -fTtt -- python manage.py test
```

pselect6(2) costs 1s, who called it?

```
...  
01:44:55.406993 pselect6(0, NULL, NULL, NULL, {tv_sec=1, tv_nsec=0}, NULL) = 0 (Timeout) <1.000610>  
...
```

-k prints the execution stack trace of the traced processes after each system call.

```
$ strace -e pselect6 -fTtt -k -- python manage.py test
```

```
01:47:36.402247 pselect6(0, NULL, NULL, NULL, {tv_sec=1, tv_nsec=0}, NULL) = 0 (Timeout) <1.001058>  
> /usr/lib/x86_64-linux-gnu/libc.so.6(__select+0xbd) [0x11b59d]  
> /usr/bin/python3.10(_Py_Gid_Converter+0x736) [0x279bb6]  
> /usr/bin/python3.10(PyObject_GenericGetAttr+0x624) [0x15c574]  
> /usr/bin/python3.10(_PyEval_EvalFrameDefault+0x613a) [0x14b34a]  
> /usr/bin/python3.10(PyFunction_Vectorcall+0x7c) [0x15d42c]  
> /usr/bin/python3.10(_PyEval_EvalFrameDefault+0x8ab) [0x145abb]  
> /usr/bin/python3.10(PyFunction_Vectorcall+0x7c) [0x15d42c]  
> /usr/bin/python3.10(_PyEval_EvalFrameDefault+0x8ab) [0x145abb]  
> /usr/bin/python3.10(PyFunction_Vectorcall+0x7c) [0x15d42c]  
...  
> /usr/bin/python3.10(_PyRun_AnyFileObject+0x43) [0x261793]  
> /usr/bin/python3.10(Py_RunMain+0x2be) [0x2542ce]  
> /usr/bin/python3.10(Py_BytesMain+0x2d) [0x22a70d]  
> /usr/lib/x86_64-linux-gnu/libc.so.6(__libc_init_first+0x90) [0x29d90]  
> /usr/lib/x86_64-linux-gnu/libc.so.6(__libc_start_main+0x80) [0x29e40]  
> /usr/bin/python3.10(_start+0x25) [0x22a605]
```

strace outputs C functions rather than Python functions.

```
kprobe: __x64_sys_pselect6 /comm == "python3.10"/  
{
```

0. eBPF hook: triggered by pselect6 syscall

```
    $task = (struct task_struct *)curtask;  
    $__ptr = (uint64)$task->stack;  
    $__ptr += THREAD_SIZE - TOP_OF_KERNEL_STACK_PADDING;  
    $pt_regs = ((struct pt_regs *)$__ptr) - 1;  
    $rsp = uptr($pt_regs->sp);
```

1. Get user space register

```
    $frame = uptr((struct PyFrameObject *)0);  
    $i = (uint64)0;  
    while ($i <= 200) {  
        $frame = *uptr((struct PyFrameObject**)( $rsp + 8*$i));  
        if (str($frame->ob_base.ob_base.ob_type->tp_name, 5) == "frame") {  
            break;  
        }  
        $frame = (struct PyFrameObject *)0;  
        $i += 1;  
    }  
  
    if ($frame == 0) {  
        return;  
    }
```

2. Search stack to find the PyFrameObject pointer

```
    $i = 0;  
    while ($i < 20) {  
        printf("%s:%s\n", $frame->f_code->co_filename->buf, $frame->f_code->co_name->buf);  
        $i += 1;  
        $frame = $frame->f_back;  
        if ($frame == 0) {  
            return;  
        }  
    }
```

3. Print py function name and filename

4. Unwind python stack

```
}
```

```
$ sudo bpftrace --include ../headers/cpython310.h --include ../headers/ubuntu2204.h cpython_syscall.bt
```

Attaching 1 probe...

```
/home/liangzc/.local/lib/python3.10/site-packages/django/db/backends/sqlite3/creation.py:_destroy_test_db  
/home/liangzc/.local/lib/python3.10/site-packages/django/db/backends/base/creation.py:destroy_test_db  
/home/liangzc/.local/lib/python3.10/site-packages/django/test/utils.py:teardown_databases  
/home/liangzc/.local/lib/python3.10/site-packages/django/test/runner.py:teardown_databases  
/home/liangzc/.local/lib/python3.10/site-packages/django/test/runner.py:run_tests  
/home/liangzc/.local/lib/python3.10/site-packages/django/core/management/commands/test.py:handle  
/home/liangzc/.local/lib/python3.10/site-packages/django/core/management/base.py:execute  
/home/liangzc/.local/lib/python3.10/site-packages/django/core/management/base.py:run_from_argv  
/home/liangzc/.local/lib/python3.10/site-packages/django/core/management/commands/test.py:run_from_argv  
/home/liangzc/.local/lib/python3.10/site-packages/django/core/management/__init__.py:execute  
/home/liangzc/.local/lib/python3.10/site-packages/django/core/management/__init__.py:execute_from_command_line  
/home/liangzc/src/github.com/jschwinger233/pycon2024/django27690/manage.py:main  
/home/liangzc/src/github.com/jschwinger233/pycon2024/django27690/manage.py:<module>
```

```
def _destroy_test_db(self, test_database_name, verbosity=1):  
    ...  
    cursor = self._maindb_connection.cursor()  
    time.sleep(1) # To avoid "database is being accessed by other users" errors.  
    if self._test_user_create():  
        ...
```

Root cause: `sleep(1)` in `_destroy_test_db()`

# Off-CPU profile

Off-CPU profiling explains where a process goes off-CPU + how long been off-CPU.

```
kprobe:finish_task_switch.isra.0
{
    $prev = (struct task_struct *)arg0;
    if ($prev->tgid == $1) {
        @start[$prev->pid] = nsecs;
    }

    $last = @start[tid];
    if ($last != 0) {
        // Calculate user space $rsp
        ...

        // Search PyFrameObject *
        ...

        // Unwind stack, print function and filename
        ...

        printf("offcpu time: %lld\n", nsecs - $last);

        delete(@start[tid]);
    }
}
```

eBPF hook: triggered by kernel task schedule

process switched from on-cpu to off-cpu

process switched from off-cpu to on-cpu

Off-CPU profiling is useful for investigating “process can’t utilize 100% CPU”.



#27690 closed Cleanup/optimization (fixed)

Opened 8 years ago

Closed 8 years ago

## remove sleep before dropping test db?

Reported by:	David Szotten	Owned by:	nobody
Component:	Testing framework	Version:	dev
Severity:	Normal	Keywords:	
Cc:	Russell Keith-Magee	Triage Stage:	Accepted
Has patch:	yes	Needs documentation:	no
Needs tests:	no	Patch needs improvement:	no
Easy pickings:	no	UI/UX:	no
Pull Requests:	7796 merged, 7861 merged		

### Description

Whilst looking at why a test run (with a small number of tests) was being slow, i discovered <https://github.com/django/django/blob/master/django/db/backends/base/creation.py#L281>

which calls `sleep(1)` before dropping the (test) db " to avoid "database is being accessed by other users" errors".

it looks like this line has been there since the beginning of the test framework: <https://github.com/django/django/commit/7dce86ce0220ffb9f3f579cbd1e881e988764c9d>

is this still needed? the few things i've tried seem to work fine without the sleep, but i'm not sure what other tests/checks might help make sure

BentoML issue #4760

# Symptoms

```
$ bentoml serve service.py:TestService
```

```
$ while :; do curl localhost:3000; done &>/dev/null
```

1. Memory leak (`docker stats $container`)
2. No memory leak (`ps -p $PID -o %mem,rss,vsz`)

这河里吗?



## How “ps” and “docker stats” are implemented?

`ps -o rss,vsz`  
↓  
`/proc/$PID/status`  
↓  
user space memory

`docker stats`  
↓  
`/sys/fs/cgroup/.../memory.stat`  
↓  
also kernel space memory

Yeah it's definitely a case of kernel space memory leak triggered by user space Python code, we need to:

1. Find a hook of kernel space memory allocation.
2. Filter events that are triggered by Python process
3. Do user space Python stack unwind to see the call chain causing kmem leak.

eBPF hook: triggered when kmem is allocated

```
tracepoint:kmem:mm_page_alloc /pid == $1/  
{  
    // Calculate user space $rsp in kernel mode  
    ...  
    // Search PyFrameObject * from user space stack  
    ...  
    // Unwind Python stack, print function and filename  
    ...  
}
```

Exactly same code as last section

\$ bpftrace --include ../headers/cpython310.h --include ../headers/ubuntu2204.h cpython\_kmemleak.bt \$PID | group.py

```
/usr/local/lib/python3.10/tempfile.py:mkdtemp  
/usr/local/lib/python3.10/tempfile.py:__init__  
/usr/local/lib/python3.10/site-packages/bentoml/_internal/context.py:in_request  
/usr/local/lib/python3.10/contextlib.py:__enter__  
/usr/local/lib/python3.10/site-packages/_bentoml_impl/server/app.py:__call__  
/usr/local/lib/python3.10/asyncio/events.py:_run  
/usr/local/lib/python3.10/site-packages/click/core.py:main  
/usr/local/lib/python3.10/site-packages/click/core.py:__call__  
/usr/local/lib/python3.10/site-packages/_bentoml_impl/worker/service.py:<module>  
/usr/local/lib/python3.10/runpy.py:_run_code  
/usr/local/lib/python3.10/runpy.py:_run_module_as_main:54
```

Root cause: mktemp dir for each request



## bug: memory leak when I am using bentoml>=1.2 #4760

gusghrlrl101 opened this issue on May 29 · 28 comments · Fixed by #4775



Zheaoli commented on Jun 3



After debug, [@frostming](#) and me confirmed that this bug has been introduced into codebase in [#4337](#)

TL;DR;

In [#4337](#) , [@frostming](#) made a new feature: make a tmp directory per request and use the tmp directory to cache all necessary files during the request

```
with tempfile.TemporaryDirectory(prefix="bentoml-request-") as temp_dir:
    dir_token = request_directory.set(temp_dir)
    try:
        yield self
    finally:
        self._request_var.reset(request_token)
        self._response_var.reset(response_token)
        request_directory.reset(dir_token)
```



# Summary

1. Traditional debugging tools have limitations
2. eBPF is able to unwind CPython stack
3. Using different Linux hooks, we debug different issues
  - a. On-CPU performance issue: `sampling(perf_event)`
  - b. Off-CPU performance issue: `syscall(kprobe)`, off-CPU scheduling(`kprobe`)
  - c. Kernel space memory leak: `kmem alloc(tracepoint)`
4. Super fast, non-disruptive, zero CPython requirement, 100% accuracy
5. Other interesting topics
  - a. Collect Python function arguments/local vars
  - b. Execution flows (Python3.11- DTrace probes)
  - c. Topdown analysis (super super fun)



Thank you! Please don't ask tough questions 😬😭

